

A Unified Framework for Low-Resource African Natural Language Processing: Cross-Lingual Transfer, Embedding Augmentation, and Knowledge Distillation

^{1st} Omprakash Pugazhendhi
Dept. of Computer Science and Engineering
Vellore Institute of Technology
Chennai, India
omprakash.2021@vitstudent.ac.in

^{2nd} Vukosi Marivate
Dept. of Computer Science
University of Pretoria
Pretoria, South Africa
vukosi.marivate@up.ac.za

Abstract—Natural language processing for African languages is held back by three problems that compound each other: there is very little labelled training data, the languages are highly diverse even within the same family, and the best-performing models are far too large to run in the low-bandwidth, limited-compute environments where these languages are most spoken. This paper addresses all three barriers in a unified study spanning three experiments on publicly available African-language benchmarks.

In the first experiment, we benchmark how well six model architectures transfer knowledge from Kinyarwanda to the closely related Kirundi language for a 12-class news classification task. We introduce a *catastrophic forgetting percentage* (CF%) to measure how much each model degrades on its source language after retraining on the target. AfriBERTa achieves the best Kirundi accuracy at 88.30% with only 5.14% source-language degradation.

In the second experiment, we propose MAGE (Multi-Head Attention Guided Embeddings), a fully language-independent pipeline for tweet sentiment analysis. MAGE augments sentence embeddings with small random perturbations, refines them through an autoencoder, and classifies with a four-head attention ensemble. On combined AfriSenti data for Kinyarwanda, Swahili, and Xitsonga, MAGE with a denoising autoencoder improves accuracy by 3.64 percentage points over a logistic regression baseline.

In the third experiment, we compress AfroXMLR-Large (559.9M parameters, 2.1 GB) into AfroXMLR-Comet (68.9M parameters, 263 MB) by simultaneously training the student to match the teacher’s output probabilities and internal attention patterns. AfroXMLR-Comet achieves 14.0 ms inference latency — $21\times$ faster than the teacher — and retains 89.2% of teacher F1 across five African languages.

Index Terms—low-resource NLP, African languages, cross-lingual transfer, knowledge distillation, data augmentation, Kinyarwanda, Kirundi, sentiment analysis

I. INTRODUCTION

Africa is home to more than 2,000 living languages, yet fewer than twenty appear in mainstream NLP benchmarks [6].

Supported by the Data Science for Social Impact (DSFSI) Research Laboratory, University of Pretoria.

This gap means that hundreds of millions of speakers of languages such as Kinyarwanda, Kirundi, Swahili, Hausa, Igbo, and Yoruba are largely excluded from AI-powered services that depend on language understanding — search, content moderation, health information delivery, and educational tools among them.

Three structural barriers drive this exclusion. The first is *data scarcity*: even the largest available labelled corpora for Bantu languages contain only tens of thousands of sentences, compared with billions for English. Training a model from scratch on such data is impractical. The second is *linguistic diversity*: even closely related languages like Kinyarwanda and Kirundi share only partial vocabulary, so transferring a model across them is not straightforward. The third is *computational asymmetry*: state-of-the-art multilingual models contain hundreds of millions of parameters, making them impractical for deployment in Sub-Saharan Africa where GPU infrastructure is scarce and connectivity is intermittent.

This paper presents three experiments that directly address each barrier.

Contribution 1. We systematically benchmark zero-shot cross-lingual transfer from Kinyarwanda to Kirundi using the KINNEWS/KIRNEWS news corpus [7]. Six architectures are compared: mBERT [1], AfriBERTa [3], BantuBERTa [5], BiGRU [23], TextCNN [21], and CharCNN [22]. We define a *catastrophic forgetting percentage* (CF%) as the fraction of source-language accuracy lost after target-language fine-tuning.

Contribution 2. We implement MAGE, a three-stage pipeline for low-resource sentiment classification using AfriSenti data [6]. The pipeline applies Language-Independent Data Augmentation (LiDA) in the embedding space, refines augmented embeddings with an autoencoder (DAE or VAE), and classifies using a multi-head attention ensemble. The entire pipeline is language-agnostic: it requires no dictionaries, translation systems, or language-specific lexical resources.

Contribution 3. We implement AfroXMLR-Comet, a com-

pact student model distilled from AfroXLMR-Large [4] using a hybrid loss that combines KL divergence on soft output distributions and MSE on internal attention maps. The student is then fine-tuned on five African languages and evaluated on AfriSenti.

The paper is organised as follows. Section II reviews related work. Section III describes the datasets. Section IV details all three methodologies. Section V reports experimental results. Section VI interprets the findings. Section VII concludes.

II. RELATED WORK

A. Multilingual Pretrained Language Models

mBERT [1] and XLM-R [2] demonstrated that a single Transformer pretrained on 100+ languages can transfer zero-shot to unseen target languages through shared sub-word vocabulary. AfroXLMR [4] further adapted XLM-R via continued pretraining on 17 African languages. AfriBERTa [3] and BantuBERTa [5] are purpose-built models trained exclusively on African corpora, enabling stronger within-family transfer for Bantu languages.

B. Cross-Lingual Transfer and Catastrophic Forgetting

Zero-shot cross-lingual transfer exploits shared sub-word tokens across related languages [2]. A known failure mode is catastrophic forgetting [15]: when fine-tuned on a new language, a model can lose accuracy on the original. Elastic Weight Consolidation [16] combats this by penalising changes to the most task-critical parameters. Our CF% metric quantifies this degradation cleanly and consistently across all six architectures, without modifying the training procedure.

C. Data Augmentation for Low-Resource NLP

Easy Data Augmentation (EDA) [17] applies synonym replacement and random word operations, but requires language-specific lexical resources unavailable for most African languages. Back-translation [18] generates extra training examples via machine translation, but reliable translation systems do not exist for most Bantu languages. LiDA [9] bypasses both limitations by adding small random perturbations directly to sentence embeddings rather than manipulating text, making it fully language-agnostic. Denoising autoencoders [20] and variational autoencoders [19] then provide principled ways to regularise the perturbed representations.

D. Knowledge Distillation

Hinton et al. [11] introduced knowledge distillation: training a student to match a teacher’s soft output distributions rather than hard ground-truth labels, transferring the teacher’s confidence patterns across classes. DistilBERT [12] applied this to BERT, retaining 97% of teacher performance at 40% fewer parameters. TinyBERT [13] added intermediate layer matching including attention-map alignment. Our distillation approach follows the attention transfer idea from [14], using mean-pooled attention matrices with a learnable projection to handle the dimensional gap between a 16-head teacher and an 8-head student.

III. DATASETS

A. KINNEWS and KIRNEWS

KINNEWS and KIRNEWS [7] contain online news articles in Kinyarwanda and Kirundi across 14 topical categories (Politics, Sport, Economy, Health, Entertainment, History, Technology, Tourism, Culture, Fashion, Religion, Environment, Education, Relationships). Following [8], we remove the two smallest classes (Fashion and Tourism) and remap the remaining 12 labels to a contiguous zero-indexed range. Article titles and body text are concatenated with a [SEP] separator before tokenisation. Dataset sizes after preprocessing are shown in Table I.

TABLE I
NEWS CLASSIFICATION DATASETS (AFTER LABEL FILTERING)

Language	Train	Test	Classes
Kinyarwanda	17,014	4,254	12
Kirundi	3,690	922	12

B. AfriSenti

AfriSenti [6] provides tweet-level sentiment annotations (positive, negative, neutral) for eleven African languages from SemEval-2023 Shared Task 12. We use three languages for the MAGE augmentation pipeline (Section IV-B) and five for the distillation fine-tuning stage (Section IV-C). Class distributions are moderately imbalanced, with neutral tweets underrepresented in smaller language subsets. Dataset splits are shown in Table II.

TABLE II
AFRISENTI SENTIMENT DATASET SPLITS

Language	Experiment	Train	Test
Kinyarwanda	MAGE + Distil	3,302	1,026
Swahili	MAGE + Distil	1,810	748
Xitsonga	MAGE only	1,006	252
Hausa	Distil only	14,172	5,303
Igbo	Distil only	10,192	3,682
Yoruba	Distil only	8,522	4,515

IV. METHODOLOGY

A. Cross-Lingual Transfer Benchmark

1) *Transformer Models*: We fine-tune three pretrained transformer models on the Kinyarwanda training set and evaluate on Kirundi under two conditions: (1) zero-shot, with no Kirundi exposure at all, and (2) after further fine-tuning on the Kirundi training set. All models are loaded as 12-class sequence classifiers via HuggingFace Transformers [26].

The shared training configuration is: batch size 32, learning rate 5×10^{-5} , weight decay 0.01, 500 warm-up steps, and maximum input length 128 tokens. mBERT and BantuBERTa train for 8 epochs; AfriBERTa trains for 25 epochs to allow its

smaller, Africa-specific vocabulary to fully converge. Checkpoints are selected by best macro F1.

2) *Traditional Neural Models*: Three non-transformer architectures are evaluated alongside the transformers. They are far smaller and cheaper to train, making them useful baselines for settings with very limited compute.

BiGRU. A two-layer bidirectional GRU [23] over 50-dimensional skip-gram Word2Vec embeddings [24] trained from scratch on the combined Kinyarwanda and Kirundi corpora (window 5, min count 5, hierarchical softmax). Hidden dimension 256; dropout 0.5; 10 training epochs with Adam.

TextCNN. The architecture from [21]: parallel convolutions with filter widths $\{3, 4, 5\}$ (150 feature maps each) over the same Word2Vec embeddings, max-over-time pooling, dropout 0.5, linear classifier. Trained for 8 epochs.

CharCNN. The architecture from [22]: 68-character vocabulary, sequences padded or truncated to 1,500 characters, six convolutional layers (256 filters, kernel 7), three max-pooling layers, two fully connected layers (1,024 units). Trained with SGD at learning rate 0.01 for 30 epochs.

3) *Catastrophic Forgetting Metric*: After a model has been trained on Kinyarwanda and evaluated on Kirundi, we fine-tune it further on Kirundi and then re-evaluate it on Kinyarwanda. The catastrophic forgetting percentage measures how much Kinyarwanda accuracy was lost during that Kirundi retraining:

$$CF\% = \frac{\text{Acc}_{\text{KIN}}^{\text{before}} - \text{Acc}_{\text{KIN}}^{\text{after}}}{\text{Acc}_{\text{KIN}}^{\text{before}}} \times 100 \quad (1)$$

“Before” is Kinyarwanda accuracy after training only on Kinyarwanda; “after” is Kinyarwanda accuracy once the same model has been further fine-tuned on Kirundi. A CF% of zero means the model kept all its source-language knowledge. As a concrete example: a model that scores 85% before and 80% after Kirundi fine-tuning has $CF\% = (85 - 80)/85 \times 100 = 5.9\%$, meaning it lost about 6% of its original capability.

B. MAGE: Augmentation and Attention Ensemble

MAGE is a four-stage pipeline: choose a sentence encoder, augment embeddings with LiDA, refine with an autoencoder, and classify with a multi-head attention ensemble. Each stage can be ablated independently.

1) *Sentence Encoding and Encoder Selection*: Four candidate encoders are compared by extracting mean-pooled last-hidden-state embeddings and classifying with logistic regression on the combined AfriSenti test set. Results are shown in Table III. AfriBERTa achieves the highest F1 (0.580) and is used for all subsequent MAGE stages.

2) *Language-Independent Data Augmentation (LiDA)*: Given an AfriBERTa sentence embedding $\mathbf{e} \in \mathbb{R}^{768}$, LiDA generates an augmented copy by adding uniform random noise independently to every dimension:

$$\tilde{\mathbf{e}} = \mathbf{e} + \boldsymbol{\epsilon}, \quad \boldsymbol{\epsilon} \sim \mathcal{U}(r_{\min}, r_{\max})^{768} \quad (2)$$

with $r_{\min} = 0$ and $r_{\max} = 0.1$. The noise nudges the sentence representation into a neighbouring region of the embedding

TABLE III
SENTENCE ENCODER COMPARISON ON COMBINED AFRISENTI TEST SET

Encoder	Acc	F1	Prec	Recall
mBERT	0.483	0.425	0.464	0.455
Afro XLM-R	0.520	0.522	0.528	0.549
BantuBERTa	0.526	0.513	0.520	0.511
AfriBERTa	0.569	0.580	0.578	0.560

space while keeping it semantically close to the original. Two augmented copies are generated per original, tripling the training set. Because the perturbation is applied to pre-computed embeddings rather than to raw text, no language-specific vocabulary or grammar rules are needed: the same procedure works identically for Kinyarwanda, Swahili, and Xitsonga.

3) *Autoencoder Refinement*: Two autoencoder variants are trained to clean and regularise the LiDA-perturbed embeddings before they are passed to the classifier.

Denoising Autoencoder (DAE). An encoder $768 \rightarrow 256 \rightarrow 64 \rightarrow 32$ (LeakyReLU, BatchNorm1d, Dropout 0.2) compresses the perturbed embedding into a 32-dimensional latent code. A decoder $32 \rightarrow 64 \rightarrow 256 \rightarrow 768$ reconstructs the original clean embedding. Training minimises mean-squared reconstruction error at learning rate 10^{-3} .

Variational Autoencoder (VAE). A shared encoder layer ($768 \rightarrow 512$) branches into two parallel heads producing mean μ and log-variance $\log \sigma^2$, each of dimension 256. A latent vector is sampled via the reparameterisation trick and decoded back through $256 \rightarrow 512 \rightarrow 768$. The training objective is the ELBO:

$$\mathcal{L}_{\text{VAE}} = \mathbb{E}[\|\hat{\mathbf{e}} - \mathbf{e}\|_2^2] + \beta D_{\text{KL}}(\mathcal{N}(\mu, \sigma^2) \parallel \mathcal{N}(0, I)) \quad (3)$$

with $\beta = 1$. The first term penalises reconstruction error (the autoencoder must reproduce the original embedding faithfully). The second term penalises deviation from a standard normal, which regularises the latent space and prevents the model from simply memorising each embedding without learning general structure. Dropout 0.2 is applied throughout.

4) *Multi-Head Attention Guided Ensemble (MAGE)*: The MAGE classifier takes a pooled embedding $\mathbf{e}$ and applies $H = 4$ independent attention heads, each with a learnable context vector $\mathbf{c}_h \in \mathbb{R}^{768}$:

$$\mathbf{o}_h = \sigma(\mathbf{e}^\top \mathbf{c}_h) \cdot \mathbf{e}, \quad \mathbf{z} = \sum_{h=1}^H \mathbf{o}_h \quad (4)$$

The sigmoid acts as a gate: when the embedding aligns well with head h ’s context vector, the gate opens and that head contributes strongly to the final representation $\mathbf{z}$; when alignment is low, the head is suppressed. Summing four independently-gated copies lets the model simultaneously attend to four different semantic aspects of the sentence — for example, one head may specialise in negation signals while another attends to intensifiers. The aggregated vector $\mathbf{z}$ is then passed

through a two-layer MLP ($768 \rightarrow 384 \rightarrow 3$, ReLU, Dropout) to produce the three sentiment class logits.

C. AfroXLMR-Comet: Knowledge Distillation

1) *Teacher and Student Architecture*: The student, AfroXLMR-Comet, is a six-layer XLM-R model with hidden dimension 256, eight attention heads, and intermediate size 1,024. Table IV shows the full XLM-R family comparison, including two intermediate variants (Base and Mini) that serve as upper-bound baselines in the distillation evaluation.

TABLE IV
XLM-R MODEL FAMILY: ARCHITECTURE AND EFFICIENCY

Property	Large	Base	Mini	Comet
Layers	24	12	6	6
Hidden size	1,024	768	384	256
Attention heads	16	12	12	8
Intermediate	4,096	3,072	1,536	1,024
Parameters	559.9M	278.0M	117.6M	68.9M
Size (MB)	2,135.9	1,060.7	448.8	263.0
Inference (ms)	293.9	100.5	30.2	14.0

2) *Hybrid Distillation Loss*: Training the student with a hybrid loss that combines two complementary signals produces better results than either signal alone.

The *response-based* component uses KL divergence between softened teacher and student output distributions:

$$\mathcal{L}_{\text{distill}} = T^2 \cdot D_{\text{KL}}\left(\text{softmax}\left(\frac{\mathbf{z}_s}{T}\right) \parallel \text{softmax}\left(\frac{\mathbf{z}_t}{T}\right)\right) \quad (5)$$

where $\mathbf{z}_s$ and $\mathbf{z}_t$ are student and teacher logit vectors and $T = 2$. Dividing by T before the softmax softens the distributions: instead of the teacher placing near-all probability on the correct class, it assigns meaningful probability to related classes too, revealing subtle patterns the student can learn from. The T^2 factor rescales gradient magnitudes to be comparable to a standard cross-entropy.

The *feature-based* component aligns the student’s last-layer attention maps with the teacher’s:

$$\mathcal{L}_{\text{attn}} = \|\mathbf{W} \bar{\mathbf{A}}_s - \bar{\mathbf{A}}_t\|_2^2 \quad (6)$$

where $\bar{\mathbf{A}} \in \mathbb{R}^{L \times L}$ is the attention matrix averaged across all heads, then flattened, and $\mathbf{W}$ is a learnable linear projection mapping the flattened student attention to the teacher’s dimensional space. This forces the student to develop similar internal attention patterns — attending to the same parts of the input — rather than just producing similar output labels. The two components are combined with equal weight:

$$\mathcal{L}_{\text{total}} = \alpha \mathcal{L}_{\text{distill}} + (1 - \alpha) \mathcal{L}_{\text{attn}}, \quad \alpha = 0.5 \quad (7)$$

3) *Training Protocol*: **Stage 1 — Distillation**. The teacher is frozen. The student is trained with AdamW at $\text{lr} = 5 \times 10^{-5}$ for up to 15 epochs, batch size 8 with 2-step gradient accumulation (effective batch 16), FP16 mixed precision, temperature $T = 2$, early stopping patience 3 and threshold 0.01. The

distillation corpus uses up to 250,000 sequences per language (Table V).

Stage 2 — Fine-tuning. The distilled student is fine-tuned independently on each AfriSenti language at $\text{lr} = 2 \times 10^{-5}$ for 5 epochs. Separating distillation from fine-tuning prevents task-specific adaptation from interfering with knowledge transfer.

TABLE V
DISTILLATION PRETRAINING CORPUS (STAGE 1)

Language	Sequences Used
Kinyarwanda	226,466
Swahili	250,000 (capped from 537,847)
Hausa	173,485
Igbo	54,410
Yoruba	52,067

V. EXPERIMENTS AND RESULTS

A. Cross-Lingual Transfer

Table VI reports accuracy and weighted F1 on the Kirundi test set before and after fine-tuning. All three transformers transfer strongly: AfriBERTa’s zero-shot accuracy (74.21%) almost matches mBERT’s after fine-tuning (84.62%), illustrating the value of Africa-specific pretraining vocabulary. Among traditional models, BiGRU reaches 83.32% accuracy after fine-tuning — competitive with the transformers — but its zero-shot accuracy (24.04%) shows that it has no transferable knowledge before direct target-language exposure. CharCNN’s post fine-tuning accuracy (48.79%) confirms that character-level features alone are insufficient for Bantu morphological complexity.

TABLE VI
CROSS-LINGUAL TRANSFER RESULTS ON KIRUNDI TEST SET

Model	Zero-Shot		After Fine-Tune	
	Acc	F1	Acc	F1
AfriBERTa	0.7421	0.7474	0.8830	0.8787
BantuBERTa	0.7454	0.7375	0.8657	0.8606
mBERT	0.5872	0.5917	0.8462	0.8422
BiGRU	0.2404	0.2300	0.8332	0.8790
TextCNN	0.2190	0.2320	0.5913	0.5732
CharCNN	0.1916	0.1621	0.4879	0.4764

Table VII shows catastrophic forgetting results. The contrast is stark: mBERT loses only 3.03% of Kinyarwanda accuracy after Kirundi fine-tuning; AfriBERTa loses 5.14%. The traditional neural models lose 71–75% — their Kinyarwanda accuracy collapses to near-chance level. The most surprising result is BantuBERTa: despite being an Africa-specific model, it forgets at 74.00%, on par with the traditional models. This strongly suggests that *breadth* of multilingual pretraining, not African-language specificity alone, determines a model’s resistance to catastrophic forgetting.

TABLE VII
CATASTROPHIC FORGETTING ON SOURCE LANGUAGE (KINYARWANDA)

Model	Acc Before	Acc After	CFW
mBERT	0.7884	0.7645	3.03
AfriBERTa	0.8498	0.8061	5.14
BantuBERTa	0.8601	0.2172	74.00
BiGRU	0.8851	0.2329	73.68
TextCNN	0.8740	0.2207	74.86
CharCNN	0.6930	0.1968	71.39

Acc Before/After: Kinyarwanda accuracy before/after Kirundi fine-tuning

B. MAGE Augmentation Results

Table VIII shows results on the combined AfriSenti test set (Kinyarwanda + Swahili + Xitsonga, 1,482 examples in total). Adding DAE or VAE alone to the LSTM or logistic regression baseline provides only marginal gains — 0.2–1.1 percentage points of accuracy. The MAGE attention ensemble produces the main improvement: MAGE+DAE improves accuracy by 3.64 percentage points over the LR baseline; MAGE+VAE improves by 2.51 points. The attention mechanism contributes +1.21 points on top of DAE refinement alone. DAE consistently outperforms VAE, likely because its deterministic reconstruction target is easier to optimise reliably with only $\sim 7,940$ combined training examples.

TABLE VIII
MAGE PIPELINE RESULTS ON COMBINED AFRISENTI TEST SET

Clf.	Aug.	Acc	Prec	Recall	F1
LSTM	None	0.5680	0.5589	0.5554	0.5566
LSTM	+ DAE	0.5739	0.5659	0.5697	0.5672
LSTM	+ VAE	0.5769	0.5681	0.5562	0.5589
LR	None	0.5710	0.5636	0.5596	0.5592
LR	+ DAE	0.5729	0.5672	0.5624	0.5619
LR	+ VAE	0.5735	0.5659	0.5621	0.5618
MAGE	+ DAE	0.6074*	—	—	—
MAGE	+ VAE	0.5961*	—	—	—

*Computed: LR baseline +3.64 pp and +2.51 pp respectively [9].

C. Knowledge Distillation Results

Table IX compares weighted F1 across five languages. AfroXLMR-Comet averages 65.28% F1 against the teacher’s 73.22%, retaining 89.2% of teacher performance while being 87.7% smaller in parameter count and $21\times$ faster at inference. Hausa and Igbo show the strongest retention (71.89% and 73.82%) because their larger fine-tuning corpora (14,172 and 10,192 samples) give the student more task-specific signal. Swahili shows the weakest retention (55.30%) due to its small fine-tuning set (1,810 samples). Notably, Comet outperforms Mini on Yoruba (66.39% vs. 63.10%) despite having 41% fewer parameters, confirming that the hybrid distillation loss

provides richer initialisation than architectural compression alone.

TABLE IX
WEIGHTED F1 (%) ACROSS LANGUAGES — DISTILLATION RESULTS

Language	Teacher	Base	Mini	Comet
Kinyarwanda	68.91	62.54	60.23	58.94
Swahili	64.10	60.81	61.88	55.30
Hausa	79.30	78.27	74.29	71.89
Igbo	79.60	78.63	75.20	73.82
Yoruba	74.20	70.25	63.10	66.39
Avg	73.22	70.10	66.94	65.28

Params: Teacher 559.9M, Base 278.0M, Mini 117.6M, Comet 68.9M.

Latency: Teacher 293.9 ms, Base 100.5 ms, Mini 30.2 ms, Comet 14.0 ms.

VI. DISCUSSION

Why does AfriBERTa outperform mBERT on cross-lingual transfer? AfriBERTa was pretrained exclusively on corpora from 11 African languages, so its sub-word vocabulary overlaps far more with Kinyarwanda and Kirundi than mBERT’s vocabulary, which must cover 104 languages within the same token budget. More shared sub-word tokens mean fewer inputs are mapped to unknown-token fallbacks, and the model’s internal representations are already calibrated to Bantu morphology from pretraining. This explains both the higher zero-shot transfer (74.21% vs. 58.72%) and the lower forgetting rate.

Why does BantuBERTa forget catastrophically despite being Africa-specific? BantuBERTa was pretrained on a narrower range of Bantu corpora, giving it a smaller and more specialised vocabulary than AfriBERTa. When fine-tuned on Kirundi, its parameters shift heavily toward the Kirundi task distribution, overwriting the Kinyarwanda signal. mBERT’s broad 104-language exposure distributes internal representations across many languages, making any single fine-tuning step a small perturbation rather than a complete overwrite. Breadth of multilingual exposure matters more than African-language specificity for forgetting resilience.

Why does MAGE benefit more from DAE than VAE? The DAE solves a well-defined reconstruction problem: given a perturbed embedding, reproduce the clean original. With only $\sim 7,940$ training examples, this converges reliably. The VAE must simultaneously learn a posterior distribution over the latent space and a reconstruction function — a harder optimisation on small data. The VAE’s stochastic sampling also introduces additional variance at test time, which hurts precision.

Why does Comet outperform Mini on Yoruba but not Swahili? For Yoruba (8,522 fine-tuning samples), the distillation loss gives Comet a richer initialisation than Mini’s random start, and the larger dataset is enough for Comet to realise that advantage. For Swahili (1,810 samples), Mini’s larger capacity (117.6M vs. 68.9M parameters) outweighs

the distillation benefit: there is not enough fine-tuning data for Comet to translate its better initialisation into higher downstream accuracy.

Practical deployment guide. The results point to a clear deployment path under different resource constraints: (1) use AfriBERTa when accuracy matters most and a 126M-parameter model can be served; (2) use AfroXLMR-Comet when latency (14 ms) or model size (263 MB) is the binding constraint; (3) use MAGE augmentation when the labelled dataset is very small and no GPU is available for transformer fine-tuning.

VII. CONCLUSION

We have presented a unified study of three complementary approaches to low-resource African NLP. Across benchmarks spanning news classification and tweet sentiment analysis in six African languages, our findings are:

- AfriBERTa achieves 88.30% accuracy on Kirundi (zero-shot: 74.21%) and loses only 5.14% of Kinyarwanda accuracy after target fine-tuning, making it the best-performing and most forgetting-resilient model for Kinyarwanda–Kirundi transfer.
- The CF% metric reveals an unexpected pattern: BantuBERTa (Africa-specific) forgets at 74.00%, matching the traditional neural models (71–75%), while mBERT (generic multilingual) retains knowledge almost perfectly (3.03%). Breadth of multilingual pretraining — not African-language specificity — determines forgetting resilience.
- MAGE with a denoising autoencoder improves combined AfriSenti accuracy by 3.64 percentage points over a logistic regression baseline, using no language-specific resources whatsoever. The attention ensemble contributes +1.21 points of that improvement beyond the augmented embeddings alone.
- AfroXLMR-Comet compresses AfroXLMR-Large by 87.7% in parameter count and runs $21\times$ faster, retaining 89.2% of teacher F1 across five languages. It outperforms the larger AfroXLMR-Mini on Yoruba, confirming that hybrid distillation provides richer initialisation than architectural compression alone.

Future work will extend the CF% analysis to named entity recognition on MasakhaNER [27], combine MAGE augmentation with distillation for very-low-resource settings, and evaluate AfroXLMR-Comet on CPU-only edge hardware representative of Sub-Saharan Africa’s infrastructure.

ACKNOWLEDGMENT

This work is conducted in collaboration with the Data Science for Social Impact (DSFSI) Research Laboratory at the University of Pretoria. We thank the Masakhane NLP community for open-sourcing African language datasets and pretrained models, and all dataset contributors whose annotation work made this research possible.

REFERENCES

- [1] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Proc. NAACL-HLT*, 2019, pp. 4171–4186.
- [2] A. Conneau et al., “Unsupervised cross-lingual representation learning at scale,” in *Proc. ACL*, 2020, pp. 8440–8451.
- [3] K. Ogueji, Y. Zhu, and J. Lin, “Small data? No problem! Exploring the viability of pretrained multilingual language models for low-resourced languages,” in *Proc. MultilingualNLP Workshop, ACL*, 2021.
- [4] J. O. Alabi, D. Adelani, P. Arnaud, and V. Marivate, “Adapting pretrained language models to African languages via multilingual adaptive fine-tuning,” in *Proc. COLING*, 2022.
- [5] B. F. P. Dossou et al., “BantuBERTa: Leveraging multilingual pretrained language models for Bantu languages,” in *Proc. AfricanNLP Workshop, EACL*, 2022.
- [6] S. H. Muhammad et al., “AfriSenti: A Twitter sentiment analysis benchmark for African languages,” in *Proc. EMNLP*, 2023.
- [7] R. A. Niyongabo, H. Qu, J. Li, and C. Ni, “KINNEWS and KIRNEWS: Benchmarking cross-lingual text classification for Kinyarwanda and Kirundi,” in *Proc. COLING*, 2020, pp. 5507–5521.
- [8] “Cross-lingual transfer of multilingual models on low resource African languages,” *arXiv:2409.10965*, 2024.
- [9] “MAGE: Multi-head attention guided embeddings for low resource sentiment classification,” *arXiv:2502.17987*, 2025.
- [10] “AfroXLMR-Comet: Multilingual knowledge distillation with attention matching for low-resource languages,” *arXiv:2502.18020*, 2025.
- [11] G. Hinton, O. Vinyals, and J. Dean, “Distilling the knowledge in a neural network,” *arXiv:1503.02531*, 2015.
- [12] V. Sanh, L. Debut, J. Chaumond, and T. Wolf, “DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter,” *arXiv:1910.01108*, 2019.
- [13] X. Jiao et al., “TinyBERT: Distilling BERT for natural language understanding,” in *Proc. EMNLP (Findings)*, 2020, pp. 4163–4174.
- [14] S. Zagoruyko and N. Komodakis, “Paying more attention to attention: Improving the performance of convolutional neural networks via attention transfer,” in *Proc. ICLR*, 2017.
- [15] M. McCloskey and N. J. Cohen, “Catastrophic interference in connectionist networks: The sequential learning problem,” *Psychol. Learning Motivation*, vol. 24, pp. 109–165, 1989.
- [16] J. Kirkpatrick et al., “Overcoming catastrophic forgetting in neural networks,” *Proc. Natl. Acad. Sci.*, vol. 114, no. 13, pp. 3521–3526, 2017.
- [17] J. Wei and K. Zou, “EDA: Easy data augmentation techniques for boosting performance on text classification tasks,” in *Proc. EMNLP-IJCNLP*, 2019, pp. 6382–6388.
- [18] R. Sennrich, B. Haddow, and A. Birch, “Improving neural machine translation models with monolingual data,” in *Proc. ACL*, 2016, pp. 86–96.
- [19] D. P. Kingma and M. Welling, “Auto-encoding variational Bayes,” in *Proc. ICLR*, 2014.
- [20] P. Vincent, H. Larochelle, Y. Bengio, and P.-A. Manzagol, “Extracting and composing robust features with denoising autoencoders,” in *Proc. ICML*, 2008, pp. 1096–1103.
- [21] Y. Kim, “Convolutional neural networks for sentence classification,” in *Proc. EMNLP*, 2014, pp. 1746–1751.
- [22] X. Zhang, J. Zhao, and Y. LeCun, “Character-level convolutional networks for text classification,” in *Proc. NeurIPS*, 2015, pp. 649–657.
- [23] K. Cho et al., “Learning phrase representations using RNN encoder-decoder for statistical machine translation,” in *Proc. EMNLP*, 2014, pp. 1724–1734.
- [24] T. Mikolov, I. Sutskever, K. Chen, G. S. Corrado, and J. Dean, “Distributed representations of words and phrases and their compositionality,” in *Proc. NeurIPS*, 2013, pp. 3111–3119.
- [25] A. Vaswani et al., “Attention is all you need,” in *Proc. NeurIPS*, 2017, pp. 5998–6008.
- [26] T. Wolf et al., “HuggingFace’s Transformers: State-of-the-art natural language processing,” in *Proc. EMNLP (System Demonstrations)*, 2020, pp. 38–45.
- [27] D. Adelani et al., “MasakhaNER: Named entity recognition for African languages,” *Trans. Assoc. Comput. Linguistics*, vol. 9, pp. 1116–1131, 2021.